


Computer Networks: Lab 1

Name: Bohdan Zasymovych

NOS: MikroTik

Topology: B

Initial configuration



Configure Alpine Hosts

Static IPs were assigned on both Alpine hosts by editing `/etc/network/interfaces` to make the changes survive reboot.

AlpineLinux-1 `/etc/network/interfaces`:

```
auto lo
iface lo inet loopback

auto eth0
iface eth0 inet static
    address 10.0.0.1
    netmask 255.255.255.0
```

AlpineLinux-2 `/etc/network/interfaces`:

```
auto lo
iface lo inet loopback

auto eth0
iface eth0 inet static
    address 10.0.0.2
    netmask 255.255.255.0
```

Configure MikroTik Switches

RouterOS doesn't bridge ports by default - each `etherX` stays isolated until added to a bridge. Same steps run on all four MikroTik VMs.

1. Check link status (R = running):

```
/interface print
```

```
Flags: R - RUNNING
Columns: NAME, TYPE, ACTUAL-MTU, MAC-ADDRESS
#  NAME      TYPE      ACTUAL-MTU  MAC-ADDRESS
0 R ether1    ether      1500        0C:B6:7B:46:00:00
1 R ether2    ether      1500        0C:B6:7B:46:00:01
2 R ether3    ether      1500        0C:B6:7B:46:00:02
3  ether4    ether      1500        0C:B6:7B:46:00:03
4 R lo        loopback   65536       00:00:00:00:00:00
```

`ether1-ether3` are up; `ether4` isn't cabled.

2. Create the bridge:

```
/interface bridge add name=bridge1
```

```
/interface bridge print
```

```
Flags: D - dynamic; X - disabled, R - running
0 R name="bridge1" mtu=auto actual-mtu=1500 l2mtu=65535 arp=enabled arp-
timeout=auto
    mac-address=0E:E4:BC:F8:27:4C protocol-mode=rstp fast-forward=yes
igmp-snooping=no
    auto-mac=yes ageing-time=5m priority=0x8000 max-message-age=20s
forward-delay=15s
```

```
transmit-hold-count=6 vlan-filtering=no dhcp-snooping=no port-cost-  
mode=long mvrp=no  
max-learned-entries=auto
```

3. Add each port to the bridge:

```
/interface bridge port add bridge=bridge1 interface=ether1  
/interface bridge port add bridge=bridge1 interface=ether2  
/interface bridge port add bridge=bridge1 interface=ether3
```

4. Verify the ports joined the bridge:

```
/interface bridge port print
```

```
Columns: INTERFACE, BRIDGE, HW, HORIZON, TRUSTED, FAST-LEAVE, BPDU-GUARD,  
EDGE, POINT-TO-POINT>
```

#	INTERFACE	BRIDGE	HW	HORIZON	TRUSTED	FAST-LEAVE	BP	EDGE	POIN
PVID	FRAME-TYPES								
0	ether1	bridge1	yes	none	no	no	no	auto	auto
1	admit-all								
1	ether2	bridge1	yes	none	no	no	no	auto	auto
1	admit-all								
2	ether3	bridge1	yes	none	no	no	no	auto	auto
1	admit-all								

After repeating this on all four MikroTik VMs, `ping 10.0.0.2` on alpine-1 and `ping 10.0.0.1` on alpine-2 both worked, with request and reply packets visible in Wireshark.

No.	Time	Source	Destination	Protocol	Length	Info
168	36.052406	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x080e, seq=48/12288, ttl=64 (reply in 169)
169	36.055399	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x080e, seq=48/12288, ttl=64 (request in 168)
170	36.583736	0c:94:96:0f:00:00	Nearest-Customer-Br...	STP	68	RST. Root = 32768/0/0c:140:60:7d:00:00 Cost = 20000 Port = 0x8001
171	36.707719	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0x7dc933a3
172	31.013865	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) request id=0x0802, seq=39/9984, ttl=64 (reply in 173)
173	31.014187	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) reply id=0x0802, seq=39/9984, ttl=64 (request in 172)
174	31.052830	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x080e, seq=49/12544, ttl=64 (reply in 175)
175	31.055457	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x080e, seq=49/12544, ttl=64 (request in 174)
176	32.015871	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) request id=0x0802, seq=40/18240, ttl=64 (reply in 177)
177	32.015484	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) reply id=0x0802, seq=40/18240, ttl=64 (request in 176)
178	32.053272	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x080e, seq=50/12800, ttl=64 (reply in 179)
179	32.056592	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x080e, seq=50/12800, ttl=64 (request in 178)
180	32.059437	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0x76e82cdd
181	32.585791	0c:94:96:0f:00:00	Nearest-Customer-Br...	STP	68	RST. Root = 32768/0/0c:140:60:7d:00:00 Cost = 20000 Port = 0x8001
182	33.014818	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) request id=0x0802, seq=41/10496, ttl=64 (reply in 183)
183	33.015101	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) reply id=0x0802, seq=41/10496, ttl=64 (request in 182)
184	33.053741	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x080e, seq=51/13056, ttl=64 (reply in 185)
185	33.057792	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x080e, seq=51/13056, ttl=64 (request in 184)
186	34.016081	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) request id=0x0802, seq=42/10752, ttl=64 (reply in 187)
187	34.016537	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) reply id=0x0802, seq=42/10752, ttl=64 (request in 186)
188	34.054213	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x080e, seq=52/13312, ttl=64 (reply in 189)
189	34.057865	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x080e, seq=52/13312, ttl=64 (request in 188)
190	34.589427	0c:94:96:0f:00:00	Nearest-Customer-Br...	STP	68	RST. Root = 32768/0/0c:140:60:7d:00:00 Cost = 20000 Port = 0x8001
191	34.672133	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0x7dc933a3
192	35.016314	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) request id=0x0802, seq=43/11008, ttl=64 (reply in 193)
193	35.016744	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) reply id=0x0802, seq=43/11008, ttl=64 (request in 192)
194	35.054807	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x080e, seq=53/13568, ttl=64 (reply in 195)
195	35.058201	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x080e, seq=53/13568, ttl=64 (request in 194)
196	36.016316	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) request id=0x0802, seq=44/11264, ttl=64 (reply in 197)
197	36.016758	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) reply id=0x0802, seq=44/11264, ttl=64 (request in 196)
198	36.055341	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x080e, seq=54/13824, ttl=64 (reply in 199)
199	36.058906	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x080e, seq=54/13824, ttl=64 (request in 198)
200	36.083929	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0x76e82cdd

Exported Configuration of the Switches

mikrotik-1:

```
[admin@mikrotik-1] > export
# 2026-09-25 15:18:47 by RouterOS 7.19.4
# system id = 2+FfVjAsZ5F
#
/interface bridge
add name=bridge1
/interface ethernet
set [ find default-name=ether1 ] disable-running-check=no
set [ find default-name=ether2 ] disable-running-check=no
set [ find default-name=ether3 ] disable-running-check=no
set [ find default-name=ether4 ] disable-running-check=no
/port
set 0 name=serial0
/interface bridge port
add bridge=bridge1 interface=ether1
add bridge=bridge1 interface=ether2
add bridge=bridge1 interface=ether3
/ip dhcp-client
# DHCP client can not run on slave or passthrough interface!
add interface=ether1
/system identity
set name=mikrotik-1
```

mikrotik-2:

```
[admin@mikrotik-2] > export
# 2026-09-25 15:28:00 by RouterOS 7.19.4
# system id = fZUCpZwANcC
#
/interface bridge
add name=bridge1
/interface ethernet
set [ find default-name=ether1 ] disable-running-check=no
set [ find default-name=ether2 ] disable-running-check=no
set [ find default-name=ether3 ] disable-running-check=no
set [ find default-name=ether4 ] disable-running-check=no
/port
set 0 name=serial0
/interface bridge port
add bridge=bridge1 interface=ether1
add bridge=bridge1 interface=ether2
add bridge=bridge1 interface=ether3
/ip dhcp-client
# DHCP client can not run on slave or passthrough interface!
add interface=ether1
/system identity
set name=mikrotik-2
```

mikrotik-3:

```
[admin@mikrotik-3] > export
# 2026-09-25 15:28:05 by RouterOS 7.19.4
# system id = fT5QtN2fq1E
#
/interface bridge
add name=bridge1
/interface ethernet
set [ find default-name=ether1 ] disable-running-check=no
set [ find default-name=ether2 ] disable-running-check=no
set [ find default-name=ether3 ] disable-running-check=no
set [ find default-name=ether4 ] disable-running-check=no
/port
set 0 name=serial0
/interface bridge port
add bridge=bridge1 interface=ether1
add bridge=bridge1 interface=ether2
add bridge=bridge1 interface=ether3
/ip dhcp-client
# DHCP client can not run on slave or passthrough interface!
add interface=ether1
/system identity
set name=mikrotik-3
```

mikrotik-4:

```
[admin@mikrotik-4] > export
# 2026-09-25 15:28:08 by RouterOS 7.19.4
# system id = AoxFsPrzEUE
#
/interface bridge
add name=bridge1
/interface ethernet
set [ find default-name=ether1 ] disable-running-check=no
set [ find default-name=ether2 ] disable-running-check=no
set [ find default-name=ether3 ] disable-running-check=no
set [ find default-name=ether4 ] disable-running-check=no
/port
set 0 name=serial0
/interface bridge port
add bridge=bridge1 interface=ether1
add bridge=bridge1 interface=ether2
add bridge=bridge1 interface=ether3
/ip dhcp-client
# DHCP client can not run on slave or passthrough interface!
add interface=ether1
/system identity
set name=mikrotik-4
```

Finding the built STP topology

To find the root bridge, run:

```
/interface bridge monitor bridge1
```

and check the **root-bridge** field - if it's **yes**, the current bridge is root.

In my case, mikrotik-4 was root:

```
[admin@mikrotik-4] > /interface bridge monitor bridge1
state: enabled
current-mac-address: 0C:40:60:7D:00:00
bridge-id: 0x8000.0C:40:60:7D:00:00
root-bridge: yes
root-bridge-id: 0x8000.0C:40:60:7D:00:00
root-path-cost: 0
root-port: none
port-count: 3
designated-port-count: 3
fast-forward: no
```

To find the role of each port:


```
/interface bridge port monitor [find interface=ether1]
```

The `role` field shows one of: `root-port`, `designated-port`, `alternate-port`. This was run for every connected port on every MikroTik VM.

Port roles (initial topology):

Switch	ether1	ether2	ether3
mikrotik-1	designated-port	designated-port	root-port
mikrotik-2	designated-port	alternate-port	root-port
mikrotik-3	designated-port	alternate-port	root-port
mikrotik-4	designated-port	designated-port	designated-port

The resulting initial topology is shown below:



On the diagram:

- **DP** – designated port
- **RP** – root port
- **AP** – alternate port
- GNS3's automatic port labels start numbering from 0, while the MikroTik CLI numbers the same interfaces starting from 1 (e.g. GNS3's `e0` = RouterOS's `ether1`).

The topology was additionally verified by running `ping 10.0.0.2` on alpine-1 and checking all inter-switch links with Wireshark. Traffic was seen on every link highlighted in the diagram (on the mikrotik-1 <-> mikrotik-2 link, only for the first ping, since that frame was unknown unicast).

Disable STP and Create a Broadcast Storm

Disable STP on a port:


```
/interface bridge port set [find interface=ether1] edge=yes
```

Note: there is no dedicated command to disable STP on just a port. This command instead marks the port as an edge port (i.e. a port connected to a host, not another switch). Since a loop is impossible on a true host-facing port, the port immediately enters forwarding mode, stops sending its own BPDUs, and ignores any it receives. Applied to an inter-switch link, this is functionally equivalent to disabling STP on that port.

Disable STP on a switch (whole bridge):

```
/interface bridge set bridge1 protocol-mode=none
```

To create a broadcast storm, I disabled STP on mikrotik-1: since it had an alternate port, disabling STP there turned that port into a forwarding port and created the loop.



This plot illustrates the number of requests per second, generated via Wireshark's **Statistics** → **I/O Graphs**. At second 140, STP was disabled on mikrotik-1, and the packet rate jumped from under 10/500ms to roughly 380 packets/500ms. Most of the flood consisted of broadcast DHCP, MNDP, and CDP traffic, which is sent by default - confirmed by inspecting the packets in Wireshark.

As an alternative, instead of disabling STP for the whole switch, I made port `ether2` (which had the `alternate-port` role) an edge port and reproduced the same broadcast storm.

Changing the Topology

In the topology that was built automatically, the root bridge was already in the correct place - only the roles of the non-root ports needed to change. This is done by adjusting port path costs.

To read a port's currently active cost, check the `actual-path-cost` field:

```
/interface bridge port monitor [find interface=ether1]
```

The default path cost on every bridge was 20000.

To reach the target topology, I needed to eliminate the links between mikrotik-3 <-> mikrotik-4 and mikrotik-2 <-> mikrotik-4. To do this, I set the path cost on ether3 of mikrotik-2 and ether3 of mikrotik-3 to 50000 - since ether3 on each is their direct link to the root (mikrotik-4). This makes that direct link more expensive than reaching the root via the two-hop path through mikrotik-1 (20000 + 20000 = 40000).

```
/interface bridge port set [find interface=ether3] path-cost=50000
```


To verify the topology change, I checked the role of every port again:

```
/interface bridge port monitor [find interface=ether1]
```

Port roles (after topology change):

Switch	ether1	ether2	ether3
mikrotik-1	designated-port	designated-port	root-port
mikrotik-2	designated-port	root-port	alternate-port
mikrotik-3	designated-port	root-port	alternate-port
mikrotik-4	designated-port	designated-port	designated-port

The resulting topology is shown below:



To check the actual path traffic takes, I used Wireshark and sent `ping -c 5 10.0.0.2` from alpine-1. Results are shown below.

The figure consists of four Wireshark packet capture screenshots arranged in a 2x2 grid, showing ICMP ping traffic between Mikrotik switches. Each screenshot displays a table of captured packets with columns: No., Time, Source, Destination, Protocol, Length, and Info.

Top-Left Screenshot (Mikrotik-3(5007) Ethernet1 to Mikrotik-1(5003) Ethernet1):

No.	Time	Source	Destination	Protocol	Length	Info
443	522.997524	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b
444	522.999700	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x081b
445	523.997690	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b
446	523.999437	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x081b
448	524.998444	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b
449	524.914246	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x081b
450	525.999930	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b
451	525.911697	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x081b
453	526.999664	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b
454	526.912113	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x081b

Top-Right Screenshot (Mikrotik-1(5003) Ethernet0 to Mikrotik-2(5005) Ethernet1):

No.	Time	Source	Destination	Protocol	Length	Info
584	516.479998	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b
585	516.481144	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x081b
588	517.480002	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b
589	517.480926	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x081b
590	518.480819	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b
591	518.485460	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x081b
593	519.481457	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b
594	519.482936	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x081b
596	520.481933	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b
597	520.483398	10.0.0.2	10.0.0.1	ICMP	98	Echo (ping) reply id=0x081b

Bottom-Left Screenshot (Mikrotik-3(5007) Ethernet2 to Mikrotik-4(5010) Ethernet2):

No.	Time	Source	Destination	Protocol	Length	Info
426	509.889474	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b, seq=0/0, ttl=64 (no response found!)

Bottom-Right Screenshot (Mikrotik-4(5010) Ethernet1 to Mikrotik-2(5005) Ethernet2):

No.	Time	Source	Destination	Protocol	Length	Info
386	481.851708	10.0.0.1	10.0.0.2	ICMP	98	Echo (ping) request id=0x081b

Full request/response traffic is visible on the mikrotik-3 <-> mikrotik-1 and mikrotik-1 <-> mikrotik-2 links. On the other links, only a single request appears - the first ping, sent as unknown unicast before the switches learned the destination MAC.

To reproduce this kind of capture in real life, one can configure port mirroring on the MikroTik switches:

```
/interface ethernet set ether1 mirror-source=ether2 mirror-target=ether1
```

and plug a laptop running Wireshark into the port traffic is mirrored to.

FDB verification: the correctness of the traffic path was additionally checked via the switches' FDBs. On mikrotik-1, -2, and -3 (all on the traffic path), entries for both Alpine hosts are present on the expected ports. On mikrotik-4, only an entry for alpine-1 is present, learned via `ether3` (the diagonal link to mikrotik-1) - this link is part of the active topology (mikrotik-1's root port), so it correctly received a copy of the initial unknown-unicast flood, even though it isn't part of the ongoing ping's forwarding path.

Picture below shows states of FDBs before and after ping on each of the MikroTik VMs

```

Columns: MAC-ADDRESS, ON-INTERFACE, BRIDGE
# MAC-ADDRESS ON-INTERFACE BRIDGE
0 D 0C:40:60:7D:00:02 ether3 bridge1
1 DL 0C:76:0A:72:00:00 bridge1 bridge1
2 DL 0C:76:0A:72:00:01 ether2 bridge1
3 DL 0C:76:0A:72:00:01 ether2 bridge1
4 D 0C:94:96:0F:00:01 ether2 bridge1
5 D 0C:B6:7B:46:00:01 ether1 bridge1
[admin@mikrotik-1] > /interface bridge host print
Flags: D - DYNAMIC; L - LOCAL
Columns: MAC-ADDRESS, ON-INTERFACE, BRIDGE
# MAC-ADDRESS ON-INTERFACE BRIDGE
0 D 0C:12:06:42:00:00 ether2 bridge1
1 D 0C:3E:1F:DC:00:00 ether1 bridge1
2 D 0C:40:60:7D:00:02 ether3 bridge1
3 DL 0C:76:0A:72:00:00 bridge1 bridge1
4 DL 0C:76:0A:72:00:01 ether2 bridge1
5 DL 0C:76:0A:72:00:02 ether3 bridge1
6 D 0C:94:96:0F:00:01 ether2 bridge1
7 D 0C:B6:7B:46:00:01 ether1 bridge1

Columns: MAC-ADDRESS, ON-INTERFACE, BRIDGE
# MAC-ADDRESS ON-INTERFACE BRIDGE
0 DL 0C:40:60:7D:00:00 ether2 bridge1
1 DL 0C:40:60:7D:00:01 ether2 bridge1
2 DL 0C:40:60:7D:00:02 ether3 bridge1
3 D 0C:76:0A:72:00:02 ether3 bridge1
4 D 0C:94:96:0F:00:01 ether3 bridge1
5 D 0C:B6:7B:46:00:01 ether3 bridge1
[admin@mikrotik-1] > /interface bridge host print
Flags: D - DYNAMIC; L - LOCAL
Columns: MAC-ADDRESS, ON-INTERFACE, BRIDGE
# MAC-ADDRESS ON-INTERFACE BRIDGE
0 DL 0C:40:60:7D:00:00 ether2 bridge1
1 DL 0C:40:60:7D:00:01 ether2 bridge1
2 DL 0C:40:60:7D:00:02 ether3 bridge1
3 D 0C:76:0A:72:00:02 ether3 bridge1
4 D 0C:94:96:0F:00:01 ether3 bridge1
5 D 0C:B6:7B:46:00:01 ether3 bridge1
[admin@mikrotik-1] > /interface bridge host print
Flags: D - DYNAMIC; L - LOCAL
Columns: MAC-ADDRESS, ON-INTERFACE, BRIDGE
# MAC-ADDRESS ON-INTERFACE BRIDGE
0 D 0C:12:06:42:00:00 ether3 bridge1
1 DL 0C:40:60:7D:00:00 bridge1 bridge1
2 DL 0C:40:60:7D:00:01 ether2 bridge1
3 DL 0C:40:60:7D:00:02 ether3 bridge1
4 D 0C:76:0A:72:00:02 ether3 bridge1
5 D 0C:94:96:0F:00:01 ether3 bridge1
6 D 0C:B6:7B:46:00:01 ether3 bridge1

```

- Alpine-1 MAC: `0C:12:06:42:00:00`
- Alpine-2 MAC: `0C:3E:1F:DC:00:00`

Additional Questions

- What is the difference between port cost and port priority?** Port cost is the value that gets added to the root path cost carried in the BPDU received on that port - in other words, it's how much "distance" using this specific port adds when calculating the total cost to reach the root bridge. By default this value depends on the link's speed (e.g. 20000 for the 1Gb/s link), but it can also be configured manually. Port priority, on the other hand, is only used as a tiebreaker - it comes into play when two ports offer an equal-cost path to the root, either when a switch is choosing its own root port among several equally-cheap candidates, or when two switches on the same segment are competing to be that segment's designated port. If the path costs already differ, priority is never even considered; cost alone decides the outcome.
- Is there sense to use Edge port in your topologies and why?** Yes, there is sense to use edge ports in our topologies. They should be used for the links between switches and hosts (alpine-1 <-> mikrotik-3 and alpine-2 <-> mikrotik-2). Since a loop cannot be formed using these links, it is safe to open the port for forwarding directly and skip the intermediate steps STP/RSTP normally goes through, lowering the

delay on the host connection. But edge ports should not be used on any links between switches, since it can cause a broadcast storm, as shown in the lab.

3. What in practice might require changing routes automatically chosen by the protocol?

- STP chooses a link only by cost (which by default depends on link speed), but it does not account for load, latency, or the real reliability of the cables and switches on that link. Using a custom metric can account for such factors.
- Before physically removing a switch or cable, the cost of the path going through it can be increased so the topology changes smoothly instead of a sudden change on removal.
- In big networks there may be a need to explicitly set root and paths to have a predictable layout, which is easier to document and debug if needed, rather than having a layout created almost randomly by default priorities and MAC addresses of the devices.

4. Is there sense to turn on STP if there are only single switch with multiple ports

Yes, there is sense to turn on STP. Despite the fact that a loop between multiple switches cannot be created in this case, it is possible to create a loop within a single switch by connecting two of its own ports together with a cable. Such a situation should be rare, but it is possible due to human error.